

# Opiframe Oy

## Async I/O



# Async I/O

- The traditional blocking I/O model is sufficient for many applications
- However there are applications that require following:
- Check a file descriptor for possible I/O without blocking.
- Monitor multiple file descriptors to see if I/O is possible on any of them.



# Async I/O

- There are different ways of handling both of these situations
- You can use non blocking I/O for the first case and handle possible errors.
- You can use multiple threads to handle the second case.



# Async I/O

- Non blocking I/O usually polls whether I/O is possible.
- This is usually very undesirable behavior since it leads either to long latency on I/O or much overhead on processor time.
- This demand increases with multiple file descriptors.



# Async I/O

- Using multiple processes or threads to handle multiple file descriptors uses loads of memory and other resources.
- In case of processes you most likely have to also use the cumbersome IPC methods to communicate between different file descriptors.
- The problem is of course expense and complexity.



# Async I/O

- Because of the limitations of both of these approaches some other methods have been devised.
- I/O Multiplexing allows process to simultaneously monitor multiple file descriptors.
- This is handled by either **select()** or **poll()** system call.



# Async I/O

- Signal-driven I/O is a technique where a process is warned by the kernel with a signal when I/O is possible.
- This is also historically called asynchronous I/O although there is an API for POSIX Async I/O.
- When monitoring a large number of file descriptors (hundreds+) this is preferable to **select()** or **poll()**



# Async I/O

- The epoll API is a Linux-specific feature that first appeared in kernel series 2.6.
- Like I/O Multiplexing, epoll API allows process to monitor multiple file descriptors for possible I/O.
- And like signal-driven I/O epoll provides much better performance when monitoring large numbers of file descriptors.





# Async I/O

- In effect, I/O Multiplexing, signal-driven I/O and epoll are all methods to achieve same result:
- To monitor several file descriptors simultaneously to see whether any or all of them are *ready* to perform I/O.
- To be precise, *ready*, here means that the I/O system call (commonly **read()** or **write()**) would not block.



# Async I/O

- The transition of a file descriptor into a ready state is triggered by some kind of I/O event.
- These include but are not limited to arrival of input, completion of socket connection or availability of space in previously full socket send buffer.
- Note that none of these I/O techniques actually do any I/O. They merely tell when the descriptor is ready.



# Async I/O – Level vs Edge Triggered

- There are two distinct ways for I/O events to happen. These are called readiness notifications.
- Level-Triggered notification: A file descriptor is considered ready if it is possible to perform I/O without blocking.
- Edge-Triggered notification: Notification is provided if there is I/O activity on file descriptor since it was last monitored.



# Async I/O – Level vs Edge Triggered

- I/O Multiplexing (**select()** and **poll()**) uses level-triggered notification.
- Signal-driven I/O uses edge-triggered notification.
- epoll uses both.
- There are a couple of major differences in handling of these readiness notifications.



# Async I/O – Level vs Edge Triggered

- With level-triggered notification we can check the availability of the file descriptor for I/O at any time.
- So we can perform some I/O on the descriptor and check if there is still possibility for more I/O
- In other words, since we check availability at any time there is no need to read or write as much as possible when I/O notification is triggered.



# Async I/O – Level vs Edge Triggered

- By contrast when we employ edge-triggered notification we only receive the notification when I/O event occurs.
- Also we do not usually know the amount of I/O is possible at the time of the notification.
- Therefore, all programs that employ edge-triggered notification should be designed according to following two rules:



# Async I/O – Level vs Edge Triggered

- After the notification for I/O happens, the program should at some point perform as much I/O as possible on corresponding file descriptor.”
- “At some point” is a safety instruction. If you perform extremely large I/O operations on one file descriptor when monitoring several you might starve the other descriptors.
- So choosing the best possible point is part of the design.



# Async I/O – Level vs Edge Triggered

- When dealing with looped I/O handling all your file descriptors should be in non blocking mode.
- Since we loop around a command to read from or write into a file descriptor at some point the I/O will end and the system call would otherwise block.
- So we perform these I/O commands until we get either **EAGAIN** or **EWouldBlock** error message from **read()** or **write()**.





# Async I/O – Non Blocking I/O

- Some examples of the use of non blocking I/O
- The reason already explained before
- If multiple processes (or threads) are performing I/O on the same file descriptor then from particular processes view the file descriptor might change state between the check and actual I/O handling.
- If we are writing a large enough block into stream socket even after level-triggered notification, the call will block.



# Async I/O – I/O Multiplexing

- We perform I/O multiplexing using either of two system calls which essentially do the same thing.
- The calls are **select()** and **poll()**.
- We can use these to monitor file descriptors for regular files, terminals, pseudoterminals, pipes, FIFOs, sockets and some types of character devices.
- Both allow process to block indefinitely or set a timeout.



# Async I/O – I/O Multiplexing

- The `select()` system call has following prototype

```
int select(int ndfs, fd_set *readfds, fd_set *writefds, fd_set *exceptfds, struct timeval *timeout);
```

- `select()` blocks until one or more file descriptors becomes ready or the timeout expires.
- The `fd_set` arguments specify the sets of file descriptors to monitor.
- The `ndfs` must be one greater than the largest number of file descriptors in any of the three sets



# Async I/O – I/O Multiplexing

- `readfds` is the set of file descriptors to be tested to see if input is possible
- `writefds` is the set of file descriptors to be tested to see if output is possible
- and `exceptfds` is the set of file descriptors to be tested for exceptional conditions



# Async I/O – I/O Multiplexing

- The exceptional conditions are not errors!
- There are two exceptional conditions that relate to specific file descriptor types.



# Async I/O – I/O Multiplexing

- The `fd_set` type is implemented as a bit mask.
- It is opaque and should not be handled directly.
- The macros to handle `fd_sets` are:
- **FD\_ZERO**(`fd_set *set`) for zeroing



# Async I/O – I/O Multiplexing

- **FD\_SET**(int *fd*, fd\_set \**set*) for adding into a set
- **FD\_CLR**(int *fd*, fd\_set \**set*) for removing from a set
- **int FD\_ISSET**(int *fd*, fd\_set \**set*) for checking if *fd* is in (return value 1) or not in (return value 0).
- The `fd_sets` have a maximum amount of file descriptors they can handle.



# Async I/O – I/O Multiplexing

- This is defined by a constant `FD_SETSIZE` which in Linux is set at 1024.
- There is no easy way to change this constant.
- Changing it requires changing the header `select.h`.
- If you are handling more than 1024 file descriptors using `epoll` is preferable anyway.





## Async I/O – I/O Multiplexing

- Before the call to `select()` all the `fd_set` structures must be initialized with `FD_ZERO` or set as `NULL` in the call.
- You add your file descriptors into the sets you are interested in.
- If we use a looping structure to loop around a `select` call this initialization must be done on all loops.
- This is because the initialization modifies these `fd_set` structures.



## Async I/O – I/O Multiplexing

- At least one of the file descriptors given in any of the `fd_set` arguments becomes ready or the call is interrupted by a signal handler.
- If the call is interrupted by a signal handler it will fail with **EINTR** as usual.
- Timeout argument brings in some portability issues. Portable programs should always initialize the **timeval struct** and ignore possible changes to it due signal interruption.



# Async I/O – I/O Multiplexing

- The return values of **select()** are as follows:
- Return value of -1 indicates an error. Possible errors include **EBADF** for bad file descriptor or **EINTR** for interruption by a signal handler.
- Return value of 0 means that the call timed out. In this case all returning `fd_sets` are empty.
- Positive integer is the total number of ready file descriptors in all sets. Same file descriptor in multiple sets is counted multiple times.



## Async I/O – I/O Multiplexing – poll()

- As with **select()** **poll()** monitors a set of file descriptors for given events.
- The difference is that in **select()** we give out three distinct sets of file descriptors but with **poll()** we give out one list of file descriptors and their corresponding events.
- The prototype for **poll()** is:

```
int poll(struct pollfd fds[], nfd_t nfd, int timeout);
```



# Async I/O – I/O Multiplexing – poll()

- The **struct pollfd** has three members:

```
struct pollfd {  
    int fd;  
    short events;  
    short revents;  
}
```

- **fd** is the specified file descriptor.
- **events** are the flags specifying the events we are interested in.
- **revents** are the returned event flags.



## Async I/O – I/O Multiplexing – poll()

- The input events concerned with either **events** or **revents** are:
- **POLLIN** (events/revents) data can be read
- **POLLRDNORM** (events/revents) same as **POLLIN**
- **POLLPRI**(events/revents) priority data can be read.
- **POLLRDHUP**(events/revents) shutdown on the peer socket.



## Async I/O – I/O Multiplexing – poll()

- The output events concerned with either **events** or **revents**:
- **POLLOUT** (events/revents) data can be written
- **POLLWRNORM**(events/revents) same as **POLLOUT**, used sometimes with sockets.
- **POLLWRBAND**(events/revents) priority data can be written.



## Async I/O – I/O Multiplexing – poll()

- Other flags, used with **revents**.
- **POLLERR**, an error occurred
- **POLLHUP**, a hang up occurred
- **POLLNVAL**, file descriptor is not open.





## Async I/O – I/O Multiplexing – poll()

- It is possible to define events as 0 if you are not interested in events on particular file descriptor.
- This makes it possible to temporarily not to be interested in a file descriptor without changing the whole **fds** list.



## Async I/O – I/O Multiplexing – poll()

- The **timeout** argument has three distinct behaviors.
- If the value is -1, **poll()** blocks until one of the file descriptors listed is ready or a signal is caught.
- If the value is 0, just check whether anyone is ready, do not block.
- If value is greater than 0, block for up to **timeout** milliseconds, or until one file descriptor is ready or a signal is caught.



## Async I/O – I/O Multiplexing – poll()

- Return values for **poll()** are following:
- -1 indicates an error similar to **select()**
- 0 indicates that the call timed out before any file descriptor became ready.
- Positive integer value indicates the one or more file descriptors are ready. The returned value is the number of **pollfd** structures in the fds array that have a nonzero **revents** field.



# Is a file descriptor ready?

- Correctly using **select()** or **poll()** requires understanding of what it means when a file descriptor is ready.
- The gist of the idea is here: A file descriptor (in blocking mode) is considered to be ready if a call to an I/O function would not block, ***regardless of whether the function would actually transfer data.***
- So the idea is not data transfer but whether an I/O operation would block.



# Implementation details

- Within the kernel, **select()** and **poll()** employ the same set of kernel-internal poll routines.
- These poll routines are different from **poll()** system call and should not be confused.
- Each poll routine returns information about readiness of a single file descriptor.



# Implementation details

- **poll()** system call calls the kernel poll routine for each file descriptor and places the resulting information in the corresponding **revents** field.
- For **select()** the return values of the poll routine are grouped into three distinct groups as per select calls arguments.
- For **poll()** there is the additional **POLLNVAL** which is not featured in **select()** groupings.
- For **select()** the same thing is the return value **EBADF** and -1.



# Performance

- I/O multiplexing calls **select()** and **poll()** have similar performance on most occasions.
- If the range of monitored file descriptors is small
- or if large number of file descriptors are densely packed (ie. the most or all file descriptors from 0 to some limit are being monitored)



# Performance

- Then the performance is similar.
- If a small number of file descriptors are monitored on a large range then **poll()** is superior.





# Problems

- Although these are widely used, portable and common place methods for monitoring multiple file descriptors, some problems remain.
- The CPU time required for **select()** and **poll()** increases with the number of file descriptors being monitored linearly.



# Problems

- Usually programs make repeated calls to monitor the same set of file descriptors but the kernel doesn't remember the list.
- So scaling is very bad.



# Signal Driven I/O

- Establish a signal handler for the signal to be delivered by the signal driven I/O mechanism. Default signal is SIGIO.
- Set the owner of the file descriptor using **fcntl()** system call ie.  
`fcntl(fd, F_SETOWN, pid);`
- Enable the non blocking I/O on the file descriptor.



# Signal Driven I/O

- Enable the signal-driven I/O by turning on the O\_ASYNC flag on the file status. ie:

```
flags = fcntl(fd, F_GETFL);  
fcntl(fd, F_SETFL, flags | O_ASYNC | O_NONBLOCK);
```

- The calling process can now go back and do other stuff. When I/O becomes possible the kernel generates a signal for the process and invokes the signal handler.



# Signal Driven I/O

- Since it is edge-triggered notification we need to perform as much I/O as possible.
- Assuming non blocking I/O as we do, this means looping **read()** or **write()** around until **EAGAIN** or **EWOULDBLOCK** is returned.
- Let's consider the actual requirement for signal delivery: When is I/O possible?



# Signal Driven I/O

- For terminals and pseudoterminals signal is generated whenever new input becomes available. Signal is also sent when end-of-file condition occurs
- No output possible signaling or terminal disconnect signaling is possible.



# Signal Driven I/O

- **For pipes and FIFOs read end is generated a signal when:**
- Data is written to the pipe
- The write end is closed



# Signal Driven I/O

- **The write end is generated a signal when:**
- A read from pipe increases the amount of free space so that it is possible to write PIPE\_BUF amount.
- The read end is closed.





# Signal Driven I/O

- For stream sockets signals are generated in the following circumstances:
- A new connection is received on listening socket.
- TCP **connect()** request completes.



# Signal Driven I/O

- New input is received into the socket.
- Peer closes the socket.
- Output is possible on the socket.
- Async error occurs on the socket.



# Signal Driven I/O – Refining

- In applications that need to simultaneously monitor very large numbers of file descriptors like certain network servers, signal-driven I/O provides significant performance gains against I/O multiplexing.
- To take full advantage of these we need to do two things.
- First we need to change the signal from SIGIO into something called real time signal.



## Signal Driven I/O – Refining

- Real time signals are the signals varying from SIGRTMIN to SIGRTMAX (32-64 in Linux).
- They can be queued when blocked unlike normal signals.
- First three real time signals are used by LinuxThreads.
- Programs should refer to these with SIGRTMIN+n for portability.



# Signal Driven I/O – Refining

- So employing **fcntl()** to change the signal.

```
int sig = SIGRTMIN+10;  
fcntl(fd, F_SETSIG, sig);
```

- You can also ask for a signal specified for a file descriptor.

```
sig = fcntl(fd, F_GETSIG);
```

- You need to add **#define \_GNU\_SOURCE** before any includes to make **F\_SETSIG** work to your code.



## Signal Driven I/O – Refining

- Specify **SA\_SIGINFO** flag and use the extended signal handler when using **sigaction()** to set the handler for the real time signal.
- The **siginfo\_t struct** passed on to the signal handler contains following:
  - **si\_signo**: the number of the signal.
  - **si\_fd**: the file descriptor in question.
  - **si\_code**: code indicating the type of the event.
  - **si\_band**: roughly equal to **revents** in **poll()**



# Signal Driven I/O – Refining

- Following **si\_codes** must match following **si\_bands**:
- **POLL\_IN** equals either **POLLIN** or **POLLRDNORM**
- **POLL\_OUT** equals either **POLLOUT** or **POLLWRNORM**
- **POLL\_ERR** equals **POLLERR**
- **POLL\_PRI** equals **POLLPRI** or **POLLRDNORM**
- **POLL\_HUP** equals **POLLHUP** or **POLLERR**



# epoll

- First thing to note is that epoll API is strictly Linux and has been around only since 2.6 kernel series.
- The performance of epoll scales much better than **select()** or **poll()**.
- The epoll API can be used with either level-triggered or edge-triggered notification.
- The performance of epoll and signal-driven I/O is similar but epoll is much simpler to handle.





# epoll

- Central to the epoll API is the epoll instance which is referred via an open file descriptor.
- This file descriptor is not used for I/O but to serve two purposes:
- Recording a list of file descriptors that this process has declared an interest in monitoring – the *interest* list.
- Maintaining a list of file descriptors that are ready for I/O – the *ready* list.



# epoll

- For each file descriptor monitored by epoll, we specify a bit mask indicating events that we are interested in.
- 
- These match closely with those of **poll()**.
- The membership in the ready list is a subset of membership in the interest list



# epoll

- The epoll API consists of three system calls:
- The **epoll\_create()** system call creates an instance of epoll and returns a file descriptor for reference.
- The **epoll\_ctl()** system call manipulates the interest list associated with an epoll instance.
- The **epoll\_wait()** system call returns items from the ready list associated with an epoll instance.



## epoll – creating an instance

- The **epoll\_create()** system call has following prototype:

```
int epoll_create(int size);
```

- The size argument specifies the number of file descriptors we expect to monitor.
- This argument is not the upper limit but rather a hint to the kernel about how initially dimension internal data structures.
- The return value is the epoll file descriptor or -1 if there is an error.



## epoll – modifying the interest list

- The **epoll\_ctl()** system call modifies the interest list of the epoll instance.
- The prototype for **epoll\_ctl()** is:  

```
int epoll_ctl(int epfd, int op, int fd, struct epoll_event *ev);
```
- The *epfd* argument is the epoll instance we are modifying.
- The *fd* argument is the actual file descriptor we are interested in. It might be on the interest list or we could be adding it there.



## epoll – modifying the interest list

- The *op* argument specifies the operation to be performed. There are three operations possible:
- **EPOLL\_CTL\_ADD** adds the file descriptor to the interest list. If we try add file descriptor which is already in the list, we get **EEXIST** error.
- **EPOLL\_CTL\_MOD** modifies the events setting for the file descriptor *fd* using the information pointed on by argument *\*ev*.
- **EPOLL\_CTL\_DEL** removes the specified file descriptor from the interest list.



## epoll – modifying the interest list

- The *\*ev* argument is a pointer to a *epoll\_event* structure which is defined as follows:

```
struct epoll_event {  
    uint32_t      events;  
    epoll_data_t  data;  
}
```

- The events subfield is a bit mask specifying the set of events that we are interested in. Those come when **epoll\_wait()** is discussed.



# epoll – modifying the interest list

- The data subfield is an union:

```
union epoll_data {  
    void *ptr;  
    int  fd;  
    uint32_t  u32;  
    uint64_t  u64;  
}
```

- One of these members can be used to specify information which is passed back to calling process via **epoll\_wait()**. So you could specify the file descriptor in question with *fd* subfield.





## epoll – waiting for events

- The **epoll\_wait()** system call returns information about ready file descriptors from the epoll instance referred.
- The prototype for **epoll\_wait()** is:  

```
int epoll_wait(int epfd, struct epoll_event *evlist, int maxevents, int timeout);
```
- The information about the ready file descriptors is returned in the *\*evlist* array.
- The *\*evlist* is allocated by the caller and contains *maxevents* number of elements.



## epoll – waiting for events

- Each element in the array *evlist* returns information about a single ready file descriptor.
- The *data* subfield of the *epoll\_event* structure contains the only distinguishing information about the file descriptor.
- It contains the information specified by the user when **epoll\_ctl()** was called. Use this to differentiate between your file descriptors.
- The *events* argument contains the events for which the file descriptor is ready.



## epoll – waiting for events

- The timeout argument determines the blocking behaviour of the `epoll_wait()`:
- If timeout is -1 block until an event occurs or signal is caught.
- If timeout is 0 do not block.
- If timeout is positive integer block up to timeout milliseconds, until an event occurs or a signal is caught.



## epoll – waiting for events

- The events for both `epoll_ctl()` and `epoll_wait()` are:
- **EPOLLIN** (ctl/wait), data can be read.
- **EPOLLPRI**(ctl/wait), high priority data can be read.
- **EPOLLRDHUP**(ctl/wait), shutdown of peer socket
- **EPOLLOUT**(ctl/wait), data can be written
- **EPOLLET**(ctl) employ edge-triggered notification
- **EPOLLONESHOT**(ctl) disable after one use
- **EPOLLERR/EPOLLHUP** (wait) error or hangup has occurred.



## epoll – edge-triggered mode

- In order to use edge-triggered notification with epoll following rules must be met:
- Make all file descriptors monitored non blocking
- Build the interest list using **epoll\_ctl()**. Remember **EPOLLET**.
- Handle I/O events using loop:
- Retrieve ready list.
- For each file descriptor process I/O until **EAGAIN** or **EWouldBlock**.

